

# Taichi - 12.06.2024

## Auteurs

- Miguel Jalube
- Romain Humair

## Description

Taichi est un langage de programmation intégré à python qui permet d'écrire du code hautement performant en utilisant le parallélisme et la vectorisation. Il est principalement utilisé pour la simulation physique, la rendu d'image et l'apprentissage profond. Taichi est conçu pour être facile à utiliser et à apprendre, il est également très flexible et extensible. Il est open source et est disponible sur github.

Taichi fait du code parallèle automatique, ce qui veut dire que le développeur n'a pas besoin de spécifier quoi que ce soit, juste ajouter l'annotation `@ti.kernel` à la fonction qu'il veut paralléliser.

## Installation

Pour installer taichi, il suffit d'exécuter la commande suivante:

```
$ pip install taichi -U
```

```
In [ ]: %pip install taichi
```

Requirement already satisfied: taichi in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (1.7.1)

Requirement already satisfied: numpy in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from taichi) (1.26.4)

Requirement already satisfied: colorama in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from taichi) (0.4.6)

Requirement already satisfied: dill in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from taichi) (0.3.8)

Requirement already satisfied: rich in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from taichi) (13.7.1)

Requirement already satisfied: markdown-it-py>=2.2.0 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from rich->taichi) (3.0.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from rich->taichi) (2.17.2)

Requirement already satisfied: mdurl~=0.1 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from markdown-it-py>=2.2.0->rich->taichi) (0.1.2)

Note: you may need to restart the kernel to use updated packages.

## Exemples

Les librairies ci-dessous sont nécessaire pour les exemples qui suivent, ces librairies ne sont pas nécessaire au bon fonctionnement de Taichi

```
In [ ]: %pip install requests pillow matplotlib numpy
```

Requirement already satisfied: requests in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (2.31.0)

Requirement already satisfied: pillow in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (10.3.0)

Requirement already satisfied: matplotlib in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (3.9.0)

Requirement already satisfied: numpy in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (1.26.4)

Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from requests) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from requests) (3.7)

Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from requests) (2.2.1)

Requirement already satisfied: certifi>=2017.4.17 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from requests) (2024.2.2)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from matplotlib) (1.2.1)

Requirement already satisfied: cycler>=0.10 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from matplotlib) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from matplotlib) (4.53.0)

Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from matplotlib) (1.4.5)

Requirement already satisfied: packaging>=20.0 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from matplotlib) (23.2)

Requirement already satisfied: pyparsing>=2.3.1 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from matplotlib) (3.1.2)

Requirement already satisfied: python-dateutil>=2.7 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from matplotlib) (2.8.2)

Requirement already satisfied: six>=1.5 in c:\users\humai\appdata\local\packages\pythonsoftwarefoundation.python.3.11\_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)

Note: you may need to restart the kernel to use updated packages.

## Compteur de nombre premier (exemple de Taichi)

Exemple fourni par Taichi pour illustrer comment Taichi peut améliorer la performance d'un code Python. Cet exemple compte le nombre de nombre premier inférieur à N

### Python

```
In [ ]: import time

N = 1000000

def is_prime(n: int):
    result = True
    for k in range(2, int(n**0.5) + 1):
        if n % k == 0:
            result = False
            break
    return result

def count_primes(n: int) -> int:
    count = 0
    for k in range(2, n):
        if is_prime(k):
            count += 1

    return count

start = time.perf_counter()
print(f"Number of primes: {count_primes(N)}")
time_python = time.perf_counter() - start;
print(f"time elapsed: {time_python}/s")
```

Number of primes: 78498  
time elapsed: 4.08143199999904/s

## Taichi : CPU

```
In [ ]: import taichi as ti
ti.init(arch=ti.cpu)

@ti.func
def is_prime(n: int):
    result = True
    for k in range(2, int(n**0.5) + 1):
        if n % k == 0:
            result = False
            break
    return result

@ti.kernel
def count_primes(n: int) -> int:
    count = 0
    for k in range(2, n):
        if is_prime(k):
            count += 1

    return count

start = time.perf_counter()
print(f"Number of primes: {count_primes(N)}")
time_cpu = time.perf_counter() - start
print(f"time elapsed: {time_cpu}/s")
```

[Taichi] Starting on arch=x64  
Number of primes: 78498  
time elapsed: 0.08143159999963245/s

## Taichi : GPU

```
In [ ]: import taichi as ti
ti.init(arch=ti.gpu)

@ti.func
def is_prime(n: int):
    result = True
    for k in range(2, int(n**0.5) + 1):
        if n % k == 0:
            result = False
            break
    return result

@ti.kernel
def count_primes(n: int) -> int:
    count = 0
    for k in range(2, n):
        if is_prime(k):
            count += 1

    return count

start = time.perf_counter()
print(f"Number of primes: {count_primes(N)}")
time_gpu = time.perf_counter() - start
print(f"time elapsed: {time_gpu}/s")
```

[Taichi] Starting on arch=cuda  
 Number of primes: 78498  
 time elapsed: 0.10463290000006964/s

## Taichi : Cuda

```
In [ ]: import taichi as ti
ti.init(arch=ti.cuda)

@ti.func
def is_prime(n: int):
    result = True
    for k in range(2, int(n**0.5) + 1):
        if n % k == 0:
            result = False
            break
    return result

@ti.kernel
def count_primes(n: int) -> int:
    count = 0
    for k in range(2, n):
        if is_prime(k):
            count += 1

    return count

start = time.perf_counter()
print(f"Number of primes: {count_primes(N)}")
```

```
time_cuda = time.perf_counter() - start
print(f"time elapsed: {time_cuda}/s")
```

[Taichi] Starting on arch=cuda  
 Number of primes: 78498  
 time elapsed: 0.09687579999990703/s

## Taichi : Vulkan

```
In [ ]: import taichi as ti
        ti.init(arch=ti.vulkan)

        @ti.func
        def is_prime(n: int):
            result = True
            for k in range(2, int(n**0.5) + 1):
                if n % k == 0:
                    result = False
                    break
            return result

        @ti.kernel
        def count_primes(n: int) -> int:
            count = 0
            for k in range(2, n):
                if is_prime(k):
                    count += 1

            return count

        start = time.perf_counter()
        print(f"Number of primes: {count_primes(N)}")
        time_vulkan = time.perf_counter() - start
        print(f"time elapsed: {time_vulkan}/s")
```

[Taichi] Starting on arch=vulkan  
 Number of primes: 78498  
 time elapsed: 0.04880860000048415/s

## Comparaison

```
In [ ]: import matplotlib.pyplot as plt

        # Execution times
        execution_times = [time_python, time_cpu, time_gpu, time_cuda, time_vulkan]

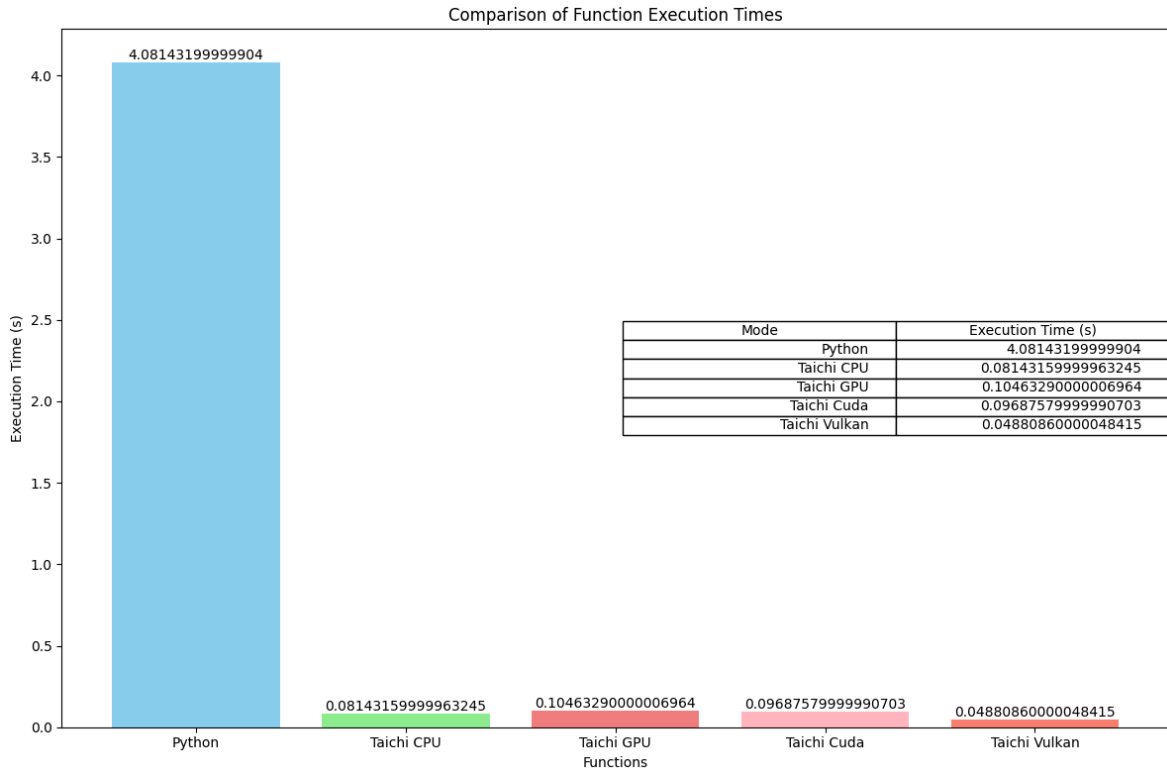
        function_names = ['Python', 'Taichi CPU', 'Taichi GPU', 'Taichi Cuda', 'Taichi Vulkan']
        colors = ['skyblue', 'lightgreen', 'lightcoral', 'lightpink', 'salmon']

        # bar chart
        # bar chart
        fig, ax = plt.subplots(figsize=(12, 8))
        bars = ax.bar(function_names, execution_times, color=colors)
        ax.set_xlabel('Functions')
        ax.set_ylabel('Execution Time (s)')
        ax.set_title('Comparison of Function Execution Times')

        for bar, time in zip(bars, execution_times):
            height = bar.get_height()
            ax.text(bar.get_x() + bar.get_width() / 2, height, f'{time}', ha='center', v
```

```
plt.subplot(1, 2, 2)
plt.axis('off') # Hide axes
table_data = list(zip(function_names, execution_times))
plt.table(cellText=table_data, colLabels=['Mode', 'Execution Time (s)'], loc='center')
plt.tight_layout()

plt.show()
```



On peut observer que le temps d'exécution est significativement réduit avec Taichi

Même observation mais en ne gardant que les backends Taichi

```
In [ ]: import matplotlib.pyplot as plt

# Execution times
execution_times = [time_cpu, time_gpu, time_cuda, time_vulkan]

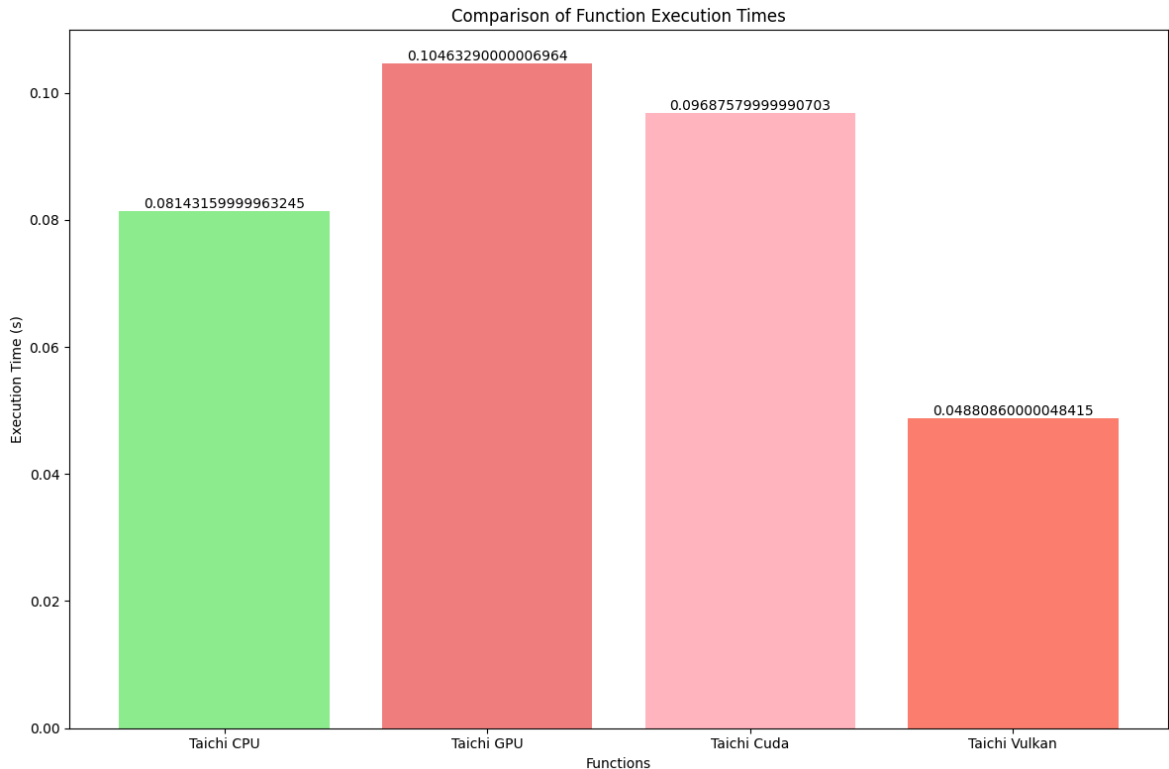
function_names = ['Taichi CPU', 'Taichi GPU', 'Taichi Cuda', 'Taichi Vulkan']
colors = ['lightgreen', 'lightcoral', 'lightpink', 'salmon']

# bar chart
fig, ax = plt.subplots(figsize=(12, 8))
bars = ax.bar(function_names, execution_times, color=colors)
ax.set_xlabel('Functions')
ax.set_ylabel('Execution Time (s)')
ax.set_title('Comparison of Function Execution Times')

for bar, time in zip(bars, execution_times):
    height = bar.get_height()
    ax.text(bar.get_x() + bar.get_width() / 2, height, f'{time}', ha='center', v

plt.tight_layout()
```

```
plt.show()
```



Pour ce cas applicatif, le backend Vulkan est plus performant. Les backends GPU et Cuda utilise la même architecture

## Calculs de matrices

### Python

```
In [ ]: import random
import time

def matrix_multiplication():
    for i in range(N):
        for j in range(N):
            temp = 0.0
            for k in range(N):
                temp += A[i][k] * B[k][j]
            result[i][j] = temp

# Définir la taille des matrices
N = 500

# Générer des matrices aléatoires
A = [[random.random() for _ in range(N)] for _ in range(N)]
B = [[random.random() for _ in range(N)] for _ in range(N)]
result = [[0 for _ in range(N)] for _ in range(N)]

# Effectuer la multiplication des matrices
start = time.perf_counter()
matrix_multiplication()
```



```
time_python2 = time.perf_counter() - start

# for row in result[:5]:
#     print(row[:5])

print(f"time elapsed: {time_python2}/s")
```

time elapsed: 14.123857399999906/s

## Taichi : GPU

```
In [ ]: import taichi as ti
import time

ti.init(arch=ti.gpu)

@ti.kernel
def matrix_multiplication():
    for i in range(N):
        for j in range(N):
            temp = 0.0
            for k in range(N):
                temp += A[i, k] * B[k, j]
            result[i, j] = temp

# Définir la taille des matrices
N = 500

# Définir les matrices en tant que champs Taichi
A = ti.field(ti.f32, shape=(N, N))
B = ti.field(ti.f32, shape=(N, N))
result = ti.field(ti.f32, shape=(N, N))

# Remplir les matrices avec des valeurs aléatoires
A.from_numpy(np.random.rand(N, N).astype(np.float32))
B.from_numpy(np.random.rand(N, N).astype(np.float32))

# Effectuer la multiplication des matrices avec Taichi
start = time.perf_counter()
matrix_multiplication()
time_taichi = time.perf_counter() - start

# print(result.to_numpy())

print(f"time elapsed: {time_taichi}/s")
```

[Taichi] Starting on arch=cuda  
time elapsed: 0.07734379999965313/s

## Python + Numpy

```
In [ ]: import numpy as np
import time

N = 500

# Définir deux matrices de taille plus importante
A = np.random.rand(N, N)
B = np.random.rand(N, N)
```

```

print(A)

# Multiplication de matrices avec NumPy
start = time.perf_counter()
result = np.dot(A, B)
time_numpy = time.perf_counter() - start

# print(result)

print(f"time elapsed: {time_numpy}/s")

```

```

[[0.9524906  0.24293516 0.97852597 ... 0.09861398 0.35515075 0.99541294]
 [0.6830298  0.63281913 0.0118799 ... 0.17070792 0.48118077 0.80409534]
 [0.85134215 0.70469624 0.1059937 ... 0.81064557 0.25958534 0.35278382]
 ...
 [0.64857507 0.80519082 0.83477724 ... 0.54586227 0.54951771 0.60987288]
 [0.24436528 0.47141184 0.5056125 ... 0.35094093 0.95560906 0.8831467 ]
 [0.120994   0.41617938 0.03950257 ... 0.86336342 0.53029776 0.86404379]]
time elapsed: 0.0030559999995602993/s

```

## Comparaison

```

In [ ]: import matplotlib.pyplot as plt

# Execution times
execution_times = [time_python2, time_taichi, time_numpy]

function_names = ['Python', 'Taichi GPU', 'Numpy']
colors = ['skyblue', 'lightgreen', 'lightcoral']

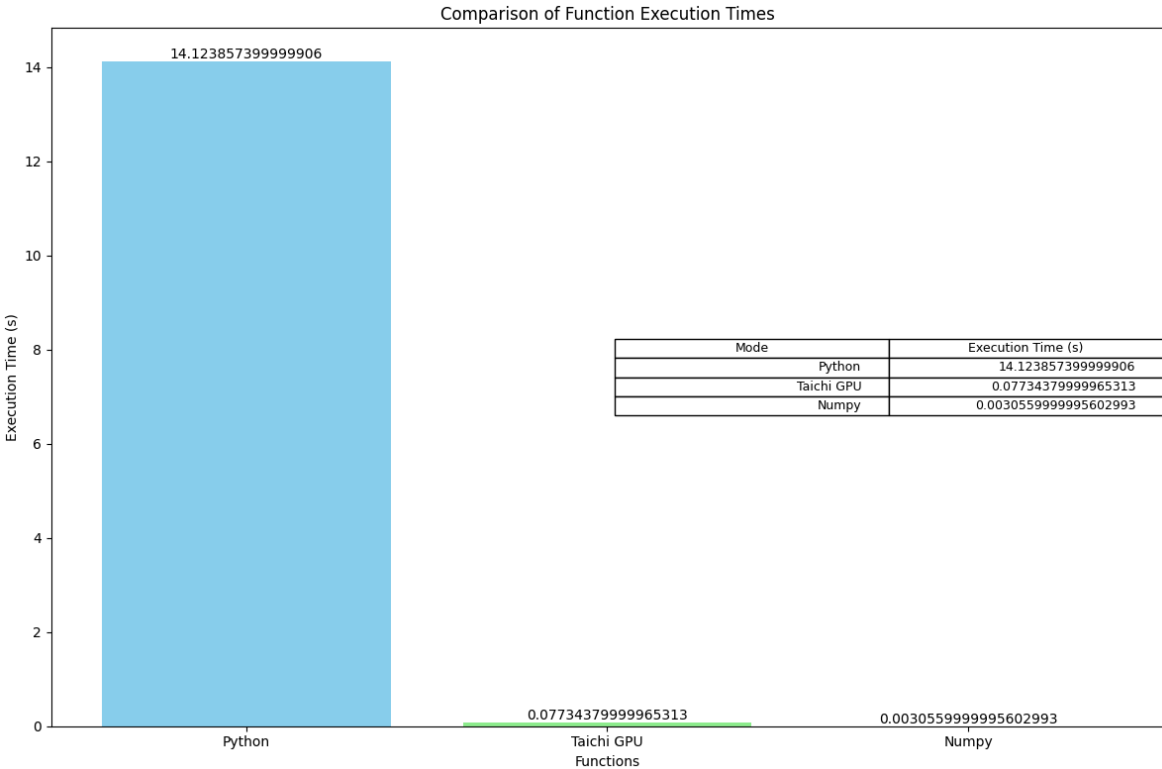
# bar chart
# bar chart
fig, ax = plt.subplots(figsize=(12, 8))
bars = ax.bar(function_names, execution_times, color=colors)
ax.set_xlabel('Functions')
ax.set_ylabel('Execution Time (s)')
ax.set_title('Comparison of Function Execution Times')

for bar, time in zip(bars, execution_times):
    height = bar.get_height()
    ax.text(bar.get_x() + bar.get_width() / 2, height, f'{time}', ha='center', v

plt.subplot(1, 2, 2)
plt.axis('off') # Hide axes
table_data = list(zip(function_names, execution_times))
plt.table(cellText=table_data, colLabels=['Mode', 'Execution Time (s)'], loc='ce
plt.tight_layout()

plt.show()

```



## Traitement d'image

Comparaison entre les performances de taichi et python sur un traitement d'images en appliquant un filtre de convolution sur une image.

```
In [ ]: import time
import numpy as np
import requests
from PIL import Image
import matplotlib.pyplot as plt
from io import BytesIO

# variables globales

# L'image est en 4k
url = "https://images7.alphacoders.com/127/1272410.jpg"

def load_image_from_url(url):
    response = requests.get(url)
    img = Image.open(BytesIO(response.content))
    return img

# afficher une image depuis Le repertoire local courant
def display_image_from_local(filename):
    img = Image.open(filename)
    # sans les axes
    plt.axis('off')
    plt.imshow(np.array(img), cmap='gray')
    plt.show()

# enregistrer l'image
def save_image(image, filename):
    image.save(filename)
```

```
# Charger L'image
image = load_image_from_url(url)

# Sauvegarder L'image
save_image(image, "../images/image.jpg")

display_image_from_local("../images/image.jpg")
```



## Python grayscale

```
In [ ]: def to_grayscale(image):
# Convertir L'image en tableau numpy
img_array = np.array(image)

# Vérifier si L'image a une couche alpha (RGBA) et L'enlever si c'est le cas
if img_array.shape[2] == 4:
    img_array = img_array[:, :, :3]

# Utiliser la formule de luminance pour convertir en niveaux de gris
grayscale_array = 0.299 * img_array[:, :, 0] + 0.587 * img_array[:, :, 1] +

# Convertir le tableau numpy en image PIL
grayscale_image = Image.fromarray(grayscale_array.astype('uint8'))

return grayscale_image

# Convertir en niveaux de gris manuellement
start_time = time.time()
grayscale_image = to_grayscale(image)
time_gray_python = time.time() - start_time

save_image(grayscale_image, "../images/vanila_grayscale.jpg")

display_image_from_local("../images/vanila_grayscale.jpg")

# Temps de conversion
print(f"Temps de conversion: {time_gray_python:.6f}s")
```



Temps de conversion: 0.288999s

## Taichi grayscale

```
In [ ]: # init taichi
import taichi as ti

ti.init(arch=ti.gpu)

def to_grayscale_ti(image):
    # Convertir l'image en tableau numpy
    img_array = np.array(image)

    # Vérifier si l'image a une couche alpha (RGBA) et l'enlever si c'est le cas
    if img_array.shape[2] == 4:
        img_array = img_array[:, :, :3]

    # Dimensions de l'image
    h, w, c = img_array.shape

    # Définir les champs Taichi pour l'image couleur et l'image en niveaux de gris
    color_image = ti.Vector.field(3, dtype=ti.u8, shape=(h, w))
    grayscale_image = ti.field(dtype=ti.u8, shape=(h, w))

    # Copier les données de l'image couleur dans le champ Taichi
    color_image.from_numpy(img_array)

    # Kernel Taichi pour convertir l'image en niveaux de gris
    @ti.kernel
    def to_grayscale():
        for i, j in ti.ndrange(h, w):
            r = color_image[i, j][0]
            g = color_image[i, j][1]
            b = color_image[i, j][2]
            grey = ti.u8(0.299 * r + 0.587 * g + 0.114 * b)
            grayscale_image[i, j] = grey

    # Exécuter le kernel
    to_grayscale()

    # Convertir le champ Taichi en tableau numpy
    grayscale_array = grayscale_image.to_numpy()
```

```

# Convertir Le tableau numpy en image PIL
greyscale_image_pil = Image.fromarray(greyscale_array)

return greyscale_image_pil

# Convertir en niveaux de gris manuellement
start_time = time.time()
greyscale_image_ti = to_greyscale_ti(image)
time_gray_taichi = time.time() - start_time

# enregistrer l'image
save_image(greyscale_image_ti, "../images/taichi_greyscale.jpg")

display_image_from_local("../images/taichi_greyscale.jpg")

# Temps de conversion
print(f"Temps de conversion: {time_gray_taichi:.6f}s")

```

[Taichi] Starting on arch=cuda



Temps de conversion: 0.263000s

## Comparaison

```

In [ ]: import matplotlib.pyplot as plt

# Execution times
execution_times = [time_gray_python, time_gray_taichi]

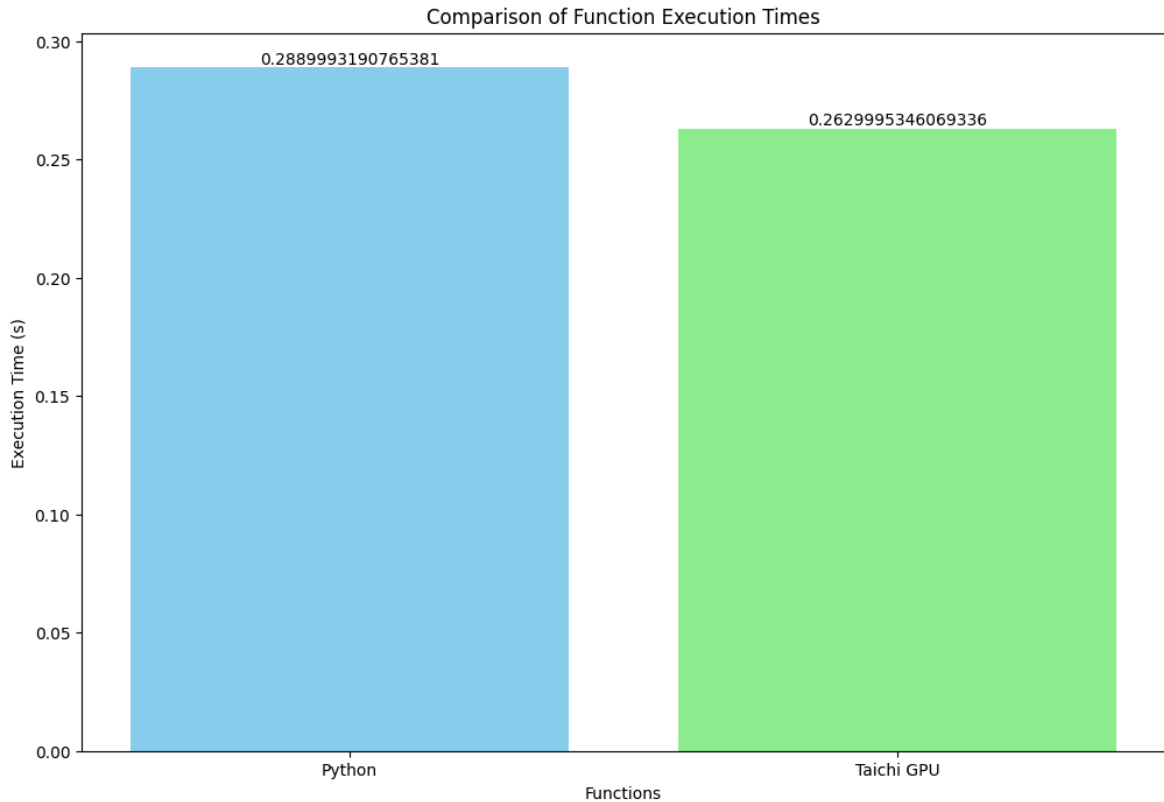
function_names = ['Python', 'Taichi GPU']
colors = ['skyblue', 'lightgreen']

# bar chart
# bar chart
fig, ax = plt.subplots(figsize=(12, 8))
bars = ax.bar(function_names, execution_times, color=colors)
ax.set_xlabel('Functions')
ax.set_ylabel('Execution Time (s)')
ax.set_title('Comparison of Function Execution Times')

for bar, time in zip(bars, execution_times):
    height = bar.get_height()

```

```
ax.text(bar.get_x() + bar.get_width() / 2, height, f'{time}', ha='center', v  
plt.show()
```



Performance équivalente entre python et taichi

## Conclusion

Selon le type d'application et le type de calcul il est possible de grandement améliorer les performances grâce à Taichi. Numpy étant basé sur des bibliothèques optimisées en C et Fortran, ce qui lui permet de bénéficier pleinement de la vectorisation, se montre plus performant que Taichi dans certains cas comme des calculs sur des tableaux.

Contrairement à ce qui est mis en avant, Taichi n'est pas si simple à prendre en main. Nous avons passé beaucoup de temps à essayer de convertir des codes pythons pour Taichi sans forcément finir par y arriver.

Taichi semble être un outil très puissant pour le calcul parallèle et la vectorisation, mais il semble avoir des problèmes avec les types de données et les calculs itératifs. Il est donc difficile de comparer les gains de performances avec python de base. L'utilisation de Taichi est plus adaptée à du traitement graphique ou calcul vectoriel.